

Goodness of Fit and Independence

STAT-S520

Arturo Valdivia

04-11-23

General Setting

- ▶ Partition the sample space of interest, S , into k events or cells
 - ▶ $E_1 \cup E_2 \cup \dots \cup E_k = S$
 - ▶ E_1, \dots, E_k are pairwise disjoint
- ▶ Test various hypotheses about the probabilities of those events.
- ▶ Given E_1, \dots, E_k , let $p_j = P(E_j)$ and the vector of cell probabilities $\vec{p} = (p_1, \dots, p_k)$
- ▶ Let Π be the set of all possible probability vectors $\vec{\pi} = (\pi_1, \dots, \pi_k)$ as long as
 - ▶ $\pi_1, \dots, \pi_k \geq 0$ and
 - ▶ $\pi_1 + \dots + \pi_k = 1$

Hypotheses

- We test

$$H_0 : \vec{p} \in \Pi_0 \quad \text{versus} \quad H_1 : \vec{p} \in \Pi_1$$

where Π_0 and Π_1 are disjoint sets of probability vectors whose union is Π .

Example 1

Construct S, E_1, \dots, E_k , and $\vec{p} = (p_1, \dots, p_k)$ under the null hypothesis that a 6-sided die is fair.

$$S = \{ \square, \square, \square, \square, \square, \square \}$$

$$E_1 = \{ \square \}, E_2 = \{ \square \}, \dots, E_6 = \{ \square \}$$

H_0 : The die is fair

$$\vec{p} = (p_1, p_2, \dots, p_6) = \left(\frac{1}{6}, \frac{1}{6}, \dots, \frac{1}{6} \right)$$

$$\vec{p} \in \Pi_0 \quad \Pi_0 = \left\{ \left(\frac{1}{6}, \frac{1}{6}, \dots, \frac{1}{6} \right) \right\}$$

$$H_0: p_1 = p_2 = p_3 = p_4 = p_5 = p_6 = \underline{\underline{\frac{1}{6}}}$$

Observed and Expected Cell Counts

- ▶ The sample: repeat the experiment n times and let o_j be the number of times that E_j appears, we call this the observed cell count of cell j .
- ▶ Goodness-of-fit tests compare observed cell counts to expected cell counts.
 - ▶ Expected cell count for cell j , e_j , is obtained assuming the null hypothesis is true.
 - ▶ If p_j is the probability of observing E_j under H_0 and the total number of observed values is n , cell j 's expected count is $e_j = p_j * n$.

Test Statistics

- ▶ Pearson's chi-squared statistic:

(greek letter)
chi-squared

$$\chi^2 = \sum_{j=1}^k \frac{(o_j - e_j)^2}{e_j}$$

- ▶ The Likelihood ratio chi-squared statistic is

$$G^2 = 2 \sum_{j=1}^k o_j \log \left(\frac{o_j}{e_j} \right)$$

- ▶ Both χ^2 and G^2 statistics can be approximated by a chi-squared distribution.

As the
chi-squared
distribution
is the sum
of independent
(std) normal squared
random variables.

← This test
is slightly
preferred.

Example 1: Fair Die (continued)

Let's assume we observed the following data (counts)

```
→ obs = c(3407, 3631, 3176, 2916, 3448, 3422)
n = sum(obs)
p = rep(1/6, 6) #probabilities under the null
exp = n*p
exp
```

```
## [1] 3333.333 3333.333 3333.333 3333.333 3333.333 3333.333
```

```
→ X2 = sum((obs - exp)^2/exp) ← Pearson Chi-squared.  
X2
```

```
## [1] 94.189
```

```
→ G2 = sum(2*obs*log(obs/exp))  
G2
```

```
## [1] 95.80227
```

Degrees of Freedom

- ▶ Under H_1 The correct degrees of freedom is the difference between the dimensions of the unrestricted and the restricted sets of possible p_1, \dots, p_k
- ▶ The unrestricted set has $k - 1$ dimensions (k probabilities, but they must sum to 1)
- ▶ The restricted set has less than $k - 1$ dimensions. It is determined by how many probabilities are free to vary.
Under H_0

Example 1 (continued)

Determine whether a 6-sided die is fair. Then

$$\underline{H_0 : p_1 = p_2 = \cdots = p_6 = \frac{1}{6}}$$

- ▶ The unrestricted set has $6 - 1 = 5$ probabilities that are free to vary.
- ▶ The null hypothesis specifies a single point, e.g.,
 $p_1 = \cdots = p_6 = 1/6$,
 - ▶ No probabilities are free to vary
 - ▶ The restricted set has dimension 0.
- ▶ The degrees of freedom needed are $df = \underbrace{(6 - 1)}_{H_1} - \underbrace{0}_{H_0} = 5$

Example 1 (continued)


```
df = (6 - 1) - 0  
1 - pchisq(X2, df)
```

```
## [1] 0
```

```
1 - pchisq(G2, df)
```

```
## [1] 0
```

↪ Reject H_0 and conclude
that the die is not fair



Simulation-Based Approach

```
die = as.character(1:6)
die.vec = rep(die, obs)
df1 = data.frame(die.vec)
null_dist <- df1 %>%
  specify(response = die.vec) %>%
  hypothesize(null = "point",
    → p = c("1" = 1/6, "2" = 1/6, "3" = 1/6, "4" = 1/6, "5" = 1/6, "6" = 1/6)
  generate(reps = 1000, type = "draw") %>%
  calculate(stat = "Chisq") ← Pearson chi-square.
null_dist %>%
  get_p_value(obs_stat = X2, direction = "greater")
  ← Pearson  $\chi^2$  for original sample.
```

Keep going
↓ a "6" = 1/6

```
## Warning: Please be cautious in reporting a p-value of 0.0000000
## approximation based on the number of 'reps' chosen in this simulation
## '?get_p_value()' for more information.
```

```
## # A tibble: 1 x 1
##   p_value = 0
```

Exercise 2 (ISI 13.4 Exercise 3)

According to Mendelian genetics, a recessive trait will appear in an offspring if and only if both parents contribute a recessive gene. If each parent has a dominant and a recessive gene, then the probability that their offspring will display the recessive trait is $1/4$.

A certain strain of tomato is either tall (dominant trait) or dwarf (recessive trait). The same strain has either cut leaves (dominant trait) or potato leaves (recessive trait). Let E_1 denote tall cut-leaf offspring, let E_2 denote tall potato-leaf offspring, let E_3 denote dwarf cut-leaf offspring, and let E_4 denote dwarf potato-leaf offspring.

$$H_0: p_1 = p_2 = p_3 = p_4 = \frac{1}{4}$$

→ Using this (see next page)

$j = 1, 2, 3, 4$

expected count

$p_i \cdot n$

Aside: This is not the only representation. For example

or (1) $H_0: p_4 = \frac{1}{4}$ $p_1 + p_2 + p_3 = \frac{3}{4}$ (but they free to change)
(2) $H_0: p_4 = \frac{1}{4}$ $p_5 = \frac{3}{4}$ where $E_5 = E_1 \cup E_2 \cup E_3$

Exercise 2 (ISI 13.4 Exercise 3 continued)

In 1931, J. W. MacArthur reported experimental results for $n = 1611$ offspring. MacArthur observed $o_1 = 926$, $o_2 = 288$, $o_3 = 293$, and $o_4 = 104$. Using this information, find:

- \vec{p} , the probability of each E_j (under H_0)
- The expected counts (under H_0)
- The test statistic
- The degrees of freedom
- The conclusion to the test

(work in R) \rightarrow check 04-13-23 lab

\hookrightarrow p-value was close to zero \rightarrow Reject H_0
Enough evidence to reject that the strain
of tomato follows the
Mendelian Genetics

Exercise 3: (ISI 13.4 Exercise 6: Using the Poisson Distribution)

Let $X(S) = \{0, 1, 2, \dots\}$. The random variable X is said to have a Poisson distribution with intensity parameter $\mu \in (0, \infty)$, if X has a probability mass function (PMF) ↗ discrete.

$$\underline{f(x)} = P(X = x) = \underline{\frac{\mu^x e^{-\mu}}{x!}}$$

We write $X \sim \text{Poisson}(\mu)$ and it can be shown that
→ $EX = \text{Var}X = \mu$. The Poisson distribution frequently arises when counting arrivals in a fixed time interval.

↪ Key result

Example 3 (ISI 13.4 Exercise 6 continued)

In 1910, E. Rutherford and M. Geiger counted the numbers of alpha-particle scintillations observed in each of $n = 2608$ 72-intervals. Now we partition $X(S)$ by setting $E_j = \{j - 1\}$ for $j = 1, \dots, 10$ and $E_{11} = \{10, 11, 12, \dots\}$. The null hypothesis states that counts of alpha-particle scintillations follow a Poisson distribution. Obtain the vector of \vec{p} that represents the null hypothesis, using the proposed partition. Estimate μ , using the following counts:

	E_1	E_2	E_3						E_{10}	E_{11}					
##	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
##	57	203	383	525	532	408	273	139	45	27	10	4	0	1	1

Observed counts.

Rule of thumb
count per cell > 5

Example 3 (ISI 13.4 Exercise 6 continued)

Using $\hat{\mu}$, find:

- The expected counts (under H_0)
- The test statistic
- The degrees of freedom
- The conclusion to the test

H_0 : Count of α -particles follows a $\text{Poisson}(\mu)$
 H_1 : It does not.

$$df_1 = 11 - 1$$

$$df_0 = 1 \leftarrow \text{Because the param. } \mu \text{ is not fixed (estimated from data)}$$

(work in R)

check 4-13-23 Lab

is not fixed (estimated from data)

Update: Using $E_j = \{j-1\}$ for $j=1, \dots, 10$
 $E_{11} = \{10, 11, 12, \dots\}$

Based on the update \rightarrow Fail to reject H_0